

# project\_understanding& System calls

- ◆ **Project Title**

# Low-Latency Linux Command Line (Linux\_CML)

High-Performance Linux CLI File Explorer in C++

---



## Overview

This project is a **high-performance Linux command-line file explorer and mini shell** implemented in **modern C++** using **POSIX/Linux system calls**.

It is engineered with a focus on:

- ✓ Clean Modular Architecture
- ✓ Direct OS interactions (no abstraction overhead)
- ✓ Low memory allocation overhead
- ✓ Deterministic, high-speed operations
- ✓ Real systems engineering practices

Tool	Purpose
<b>WSL (Ubuntu 22.04)</b>	Linux environment on Windows
<b>g++</b>	C++ compiler
<b>CMake</b>	Build system
<b>Git</b>	Version control
<b>VS Code Remote-WSL</b>	IDE with Linux environment integration
<b>perf / Valgrind</b>	Profiling and memory tools

Memory Type	Speed	Allocation	Lifetime
Stack	Very Fast	Automatic	Local Scope
Heap	Slower	Through new/malloc	Manual Control

### Example (Stack buffer)

```
char buffer[4096];
```

Fast, no malloc, ideal for predictable latency.

---

## ◆ System Call — getcwd()

C function signature:

```
char* getcwd(char* buf, size_t size);
```

Meaning:

- ✓ Ask the kernel for current working directory
- ✓ Kernel writes directory path into provided buffer
- ✓ Returns pointer on success, NULL on failure

Typical usage:

```
char buffer[PATH_MAX];
if(getcwd(buffer, sizeof(buffer)) != nullptr) {
    std::cout << buffer << "\n";
} else {
    perror("getcwd");
}
```

getcwd(buffer, sizeof(buffer))

This is core.

What is getcwd?

C library function.

Full form:

👉 get current working directory

Signature:

char\* getcwd(char\* buf, size\_t size);

Parameters

① buffer

Where to store result

② size

How big is buffer

So:

OS writes path inside buffer

🧠 Under the hood (important)

When you call:

getcwd()

Flow:

Your program

↓

libc

↓

kernel syscall

↓

kernel gets cwd

↓

copies into your buffer

↓

returns pointer

So:

Kernel → user space copy.

Low level.

---

Let's visualize memory:

Suppose path:

/home/kshivam

Memory:

buffer:

```
| / | h | o | m | e | / | k | s | h | i | v | a | m | \0 | ... |
```

Note:

\0 = null terminator

Marks end of string.

C style strings always end with \0

### "Explain this function"

"It uses a fixed-size stack buffer to avoid heap allocations and calls the POSIX getcwd system call to retrieve the current working directory. The kernel writes the path directly into the provided buffer. We check for failure using the returned pointer and use perror for error reporting. This approach is faster and more deterministic than using std::string, which is better suited for low-latency systems."



## Now quick intuition version

Think like:

```
char buffer[4096];      → empty box
getcwd()                 → OS fills box
cout                      → print box
```

## ◆ System Call — getcwd()

C function signature:

```
char* getcwd(char* buf, size_t size);
```

Meaning:

- ✓ Ask the kernel for current working directory
- ✓ Kernel writes directory path into provided buffer
- ✓ Returns pointer on success, NULL on failure

## 📌 Common Interview Questions You Can Answer

- ❓ Why use stack buffers over std::string?  
✓ Stack avoids heap, offers deterministic performance.
- ❓ What is getcwd and how does it work?  
✓ System call to retrieve current working directory into user buffer.
- ❓ Why use CMake instead of g++ directly?  
✓ Scalability, reproducible builds, multi-file support.
- ❓ What is PATH\_MAX?  
✓ Maximum path length constant in Linux.
- ❓ Why do we avoid binaries in Git?  
✓ Binaries are machine-specific and clutter repos.

# CD..Command Notes

# Linux File Explorer (CML) – cd Command Deep Notes

## 1. Problem Statement

Implement Linux shell command:

```
cd <path>
```

## System Call Used

`chdir()`

**Prototype:**

```
int chdir(const char* path);
```

**Return:**

value	meaning
0	success
-1	failure

## 4. Final Code

```
void Shell::changeDirectory(const string& path){  
    if(chdir(path.c_str()) != 0){  
        perror("cd");  
    }  
}
```

```
void Shell::changeDirectory(const string& path)
```

### Meaning

- Member function of Shell class
- Takes path as input

```
chdir(path.c_str())
```

### Problem:

chdir() expects:

⇒ const char\*

But we have:

⇒ std::string

## Solution

Use:

```
path.c_str()
```

---

## What is c\_str() ?

Returns:

```
const char*
```

pointer to internal buffer of string.

Example:

```
string path = ".."
```

Memory:

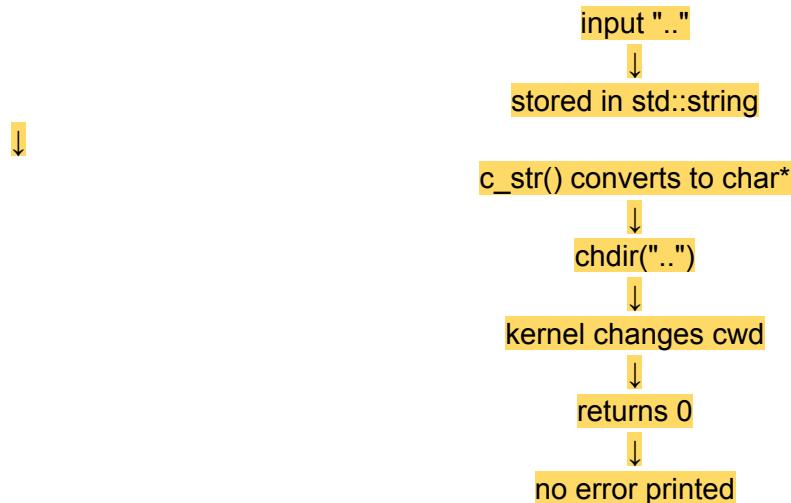
```
Heap: [ .. ' ' '\0' ]
```

c\_str() → pointer to this memory

---

## Interview point:

"POSIX APIs use C-style strings, so we convert std::string to char\* using c\_str()"



symbol	meaning
.	current
..	parent
/	root

## Q1

Why use const string& ?

→ avoids copy, faster

---

## Q2

Why c\_str() ?

→ convert to char\* for POSIX API

---

### **Q3**

Where is current directory stored?

→ kernel (process structure)

---

### **Q4**

Why not change directory manually?

→ only kernel can change cwd

---

### **Q5**

What happens on failure?

→ returns -1 and sets errno

---

### **Q6**

Stack vs heap here?

→ string on stack, char buffer on heap internally

## **Why is std::filesystem slower than chdir?**

**Answer:**

filesystem:

- allocations

- exceptions
- abstraction
- extra checks

chdir:

- direct syscall
- minimal overhead

## What happens if two threads call chdir() simultaneously?

🔥 VERY GOOD QUESTION

---

**Answer:**

cwd is **process-wide**, not thread-specific.

So:

Thread A:

`chdir(/home)`

Thread B:

`chdir(/tmp)`

Race condition.

Both affect each other.

---

### **Result:**

UNSAFE in multi-threaded programs.

## **Why is `char buffer[PATH_MAX]` preferred over `vector<char>`?**

### **Answer:**

vector:

- heap allocation
- slower
- allocator overhead

array:

- stack
  - fixed size
  - zero allocation
- 

### **HFT rule:**

Stack allocation preferred over heap to reduce latency.

## **Design question (very common)**

?

How would you design a high-performance shell for millions of file ops?

**Good answer:**

- avoid chdir frequently
- use absolute paths
- minimize syscalls
- avoid heap allocations
- use buffers
- cache metadata
- async IO if needed

# LS Command Notes

## LS Command Notes:

```
#include "include.hpp"
using namespace std;

void Shell::listDirectory(){
    DIR* dir=opendir(".");
    // . means current file. take a control. return
pointer. stored in DIR*
    if(dir==nullptr){
        perror("ls");
        return;
    }
    //internally dirent in linux
//    struct dirent {
//        ino_t d_ino;
//        char d_name[256];
//    };    Each file inside folder is represented by: dirent

    struct dirent* entry;

    while((entry=readdir(dir))!=nullptr){
        cout<<entry->d_name<<" ";
    }
    cout<<endl;
    closedir(dir);
}
```

## FIRST — Big picture (what is Is actually?)

# When you run:

1s

# What happens internally?

## Linux:

- 👉 A directory is just a file containing list of entries

## Each entry:

filename + inode number

```
DIR* dir = opendir(".");
```

Think	Type
-------	------

file	FILE*
------	-------

directory	DIR*
-----------	------

### Step 2 — What is opendir() ?

Prototype:

```
DIR* opendir(const char* path);
```

Meaning:

👉 "Kernel, open this folder and give me handle"

---

---

### Step 3 — Why "." ?

" . "

means:

👉 current directory

Same like:

`pwd`

#### Step 4 — What happens internally?

Kernel does:

1. find current directory inode

2. open it

3. allocate file descriptor

4. create DIR structure

5. return pointer

**IMPORTANT:**

`struct dirent* entry;`

## **What is dirent?**

Very important.

Linux defines:

```
struct dirent {  
    ino_t d_ino;  
    char d_name[256];  
};
```

---

### **Meaning:**

Each file inside folder is represented by:

**dirent**

---

### **Example:**

Directory:

```
src  
main.cpp  
build
```

Kernel returns:

**dirent("src")**

```
dirent("main.cpp")
dirent("build")
```

---

---

## Why pointer?

```
struct dirent*
```

Not object.

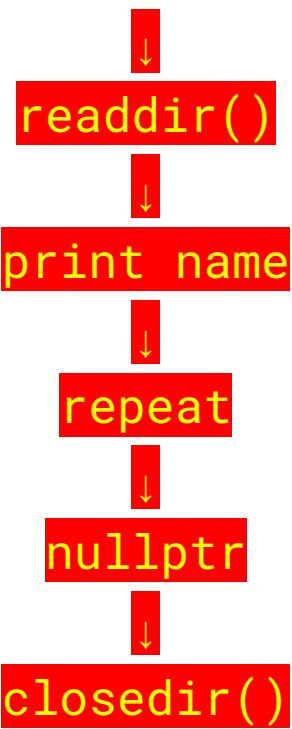
Because:

- avoid copying
- faster
- kernel gives pointer to internal buffer

Low latency design.

Flow:





SastaRivision

# 1 What is Linux (from zero)

## Definition

Linux is a **Unix-like operating system**.

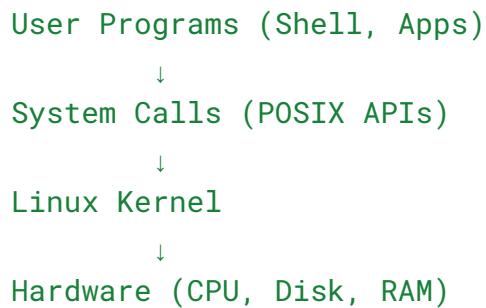
It provides:

- Process management
- Memory management
- File system
- System calls
- Device drivers

---

## Architecture

Draw this diagram in notebook:



---

## Meaning

Layer	Role
App	your shell

Syscall      `getcwd()`, `chdir()`, `fork()`

Kernel      handles real work

Hardware    executes instructions

---

## 2 What is POSIX?

### Definition

**POSIX = Portable Operating System Interface**

It is a **standard API for Unix systems**.

---

### Why needed?

Without POSIX:

- Windows → different API
- Linux → different API
- Mac → different API

Code won't be portable.

With POSIX:

Same code runs on:

- Linux
  - Mac
  - Unix
  - WSL
-

## Examples

POSIX Function	Use
getcwd()	get current dir
chdir()	change dir
opendir()	open directory
readdir()	list files
stat()	file info
fork()	create process
exec()	run program
pipe()	interprocess communication

---

## 3 What is Shell?

### Definition

Shell = **command interpreter**

It:

1. reads command
  2. parses command
  3. executes action
- 

### Example

`pwd`

Flow:

User → Shell →.getcwd() → Kernel → Path → Shell prints

---

## 4 What we built (big picture)

Draw this:

```
main()
  ↓
Shell.run()
  ↓
Read line
  ↓
Tokenize
  ↓
Match command
  ↓
Call function
```

---

## 5 Current Project Architecture

```
Linux_CML
|
├── include/shell.hpp
├── src/shell.cpp
└── src/commands/
    ├── pwd
    ├── cd
    ├── ls
    ├── stat
    ├── echo
    └── help
```

---

## Why modular?

Interview gold point:

"Each command is implemented independently to follow single responsibility and scalability principles."

---

## 6 Deep Understanding of Each Command

---

### ◆ PWD

#### Code concept

```
char buffer[PATH_MAX];
getcwd(buffer, sizeof(buffer));
```

---

#### How it works

1. buffer created on stack
  2. getcwd asks kernel
  3. kernel copies path into buffer
  4. printed
- 

#### Memory

Stack:

```
buffer[4096 bytes]
```

---

## Interview Q

Q: Why buffer needed?

A: Kernel writes result into user memory.

---

---

## ◆ CD

### Code

```
chdir(path.c_str())
```

---

### Flow

1. string → char\*
  2. kernel updates process working directory
- 

### Important

Working directory stored in:

```
task_struct (kernel)
```

---

## Interview gold

Q: Why cd cannot be separate program?

A: Because directory is property of process.

---

---

## ◆ LS

## **Code**

```
DIR* dir = opendir(".");
readdir(dir);
```

---

## **Flow**

`opendir` → open directory file

`readdir` → read entries

`closedir` → release

---

## **Memory**

`DIR*` → heap object created by libc

---

## **Interview**

Q: Why must close?

A: FD leak.

---

---

# ◆ STAT

## **Code**

```
stat(path, &info)
```

---

## **Gives**

- size
- inode

- permissions
  - type
- 

## Kernel

Reads inode metadata from filesystem.

---

## Interview

Q: stat vs lstat?

A: lstat doesn't follow symlink.

---

---

## ◆ ECHO

### Code

```
vector<string> args
```

---

### Why vector?

Modern C++ dynamic container.

Better than C arrays.

---

---

## 7 Core Engine (MOST IMPORTANT PART)

Your run() function:

## Steps

### Step 1

`getline`

---

Take full command

---

### Step 2

`stringstream`

---

Tokenize

---

### Step 3

`vector<string> tokens`

---

Store arguments

---

### Step 4

Dispatch

---

`if(command=="pwd")`

---

## Interview Question

Q: Why `getline` not `cin >>` ?

A: To support multi-argument commands.

---

---

# **8** Memory Concepts Used

You must know this for interviews:

---

## **Stack**

- char buffer[PATH\_MAX]
  - local variables
- 

## **Heap**

- vector
  - string
  - DIR\*
- 

## **Kernel memory**

- directory info
- file metadata

# **mkdir command notes**

```

PROJECTS [WSL: UBUNTU-22.04]
└─ Linux_CML
  └─ build
    └─ CMakeCache.txt
    └─ linux_cml
      └─ Makefile
    > docs
    > include
    └─ src
      └─ commands
        └─ mkdir.cpp
      > catfile.cpp
      > changeDirectory.cpp
      > copyfile.cpp
      > echo.cpp
      > executeExternal.cpp
      > history.cpp
      > listDirectory.cpp
      > touch.cpp
      > main.cpp
      > shell.cpp
    > TEST_FOR_LINKDIR
    > tests
    < .gitignore
    < CMakeLists.txt
    < README.md
  > OUTLINE
  > TIMELINE

```

```

touch.cpp
1 #include "shell.hpp"
2 #include <sys/stat.h>
3 #include <iostream>
4 #include <errno>
5
6 void Shell::makeDirectory(const
7 {
8     if(mkdir(path.c_str(), 0755
9     {
10         perror("mkdir");
11     }
12 }

```

```

Linux_CML> src> commands> mkdir.cpp
-- Configuring done
-- Generating done
-- Build files have been written to: /home/kshivam/projects/Linux_CML/build
d
● kshivam@ShivamKachhadia:~/projects/Linux_CML/build$ make
consolidate compiler generated dependencies of target linux_cml
[ 5%] Building CXX object CMakeFiles/linux_cml.dir/src/commands/redirect1
on.cpp.o
[ 11%] Linking CXX executable linux_cml
[100%] Built target linux_cml
● kshivam@ShivamKachhadia:~/projects/Linux_CML/build$ ./linux_cml
./linux_cml: command not found
● kshivam@ShivamKachhadia:~/projects/Linux_CML/build$ ./linux_cml.exe
./linux_cml.exe: command not found
● kshivam@ShivamKachhadia:~/projects/Linux_CML/build$ ./main
./main: command not found
● kshivam@ShivamKachhadia:~/projects/Linux_CML/build$ ls
CMakeCache.txt
CMakeFiles
.
linux_cml
Makefile
...
cmake_install.cmake
LINUX_CML> pwd
/home/kshivam/projects/Linux_CML/build
LINUX_CML> cd ..
LINUX_CML> pwd
/home/kshivam/projects/Linux_CML
LINUX_CML> mkdir TEST_FOR_LINKDIR
LINUX_CML>

```

## ✓ What is mkdir?

### Definition

`mkdir` = create a new directory entry in filesystem

It creates:

- a new inode
- a directory file
- adds entry in parent directory

## ✓ POSIX system call

```
int mkdir(const char* path, mode_t mode);
```

## ✓ Full code

```
#include "shell.hpp"
#include <sys/stat.h>
```

```
#include <cerrno>
#include <cstring>
#include <iostream>

void Shell::makeDirectory(const std::string& path)
{
    // call POSIX mkdir
    if (mkdir(path.c_str(), 0755) != 0)
    {
        // error handling
        std::cerr << "mkdir: " << std::strerror(errno) << "\n";
    }
}
```

Gives:

`strerror()`

Convert error number → message

---

---

## Function

```
void Shell::makeDirectory(const std::string& path)
```

Means:

- member function of `Shell`
- takes modern C++ string

---

```
if (mkdir(path.c_str(), 0755) != 0){}
```

---

`path.c_str()`

```
path.c_str()
```

WHY?

Because:

POSIX expects:

```
const char*
```

But we have:

```
std::string
```

So:

```
string → char*
```

conversion

---

---

### **mkdir call**

```
mkdir(path.c_str(), 0755)
```

Create directory with permission:

```
rwx r-x r-x
```

---

---

### **check failure**

```
!= 0
```

Because:

```
0 → success
```

-1 → failure

---

---

## error printing

`std::strerror(errno)`

Example output:

```
mkdir: File exists
mkdir: Permission denied
```

Example:

`0755`

Meaning:

```
Owner → rwx (7)
Group → r-x (5)
Others → r-x (5)
```

---

## Q1

Difference between mkdir and open?

Answer:

`open()`:

- creates regular file

`mkdir()`:

- creates directory inode
- adds ".." and "."

---

---

## Q2

Why mode 0755 not 755?

Answer:

0755 = octal  
permissions stored in octal

---

---

## Q3

What is umask?

Answer:

Default permission mask applied by kernel:

`final = mode & ~umask`

---

---

## Q4

What happens internally when mkdir called?

Answer:

- inode allocation
  - directory entry creation
  - update parent directory
  - permission set
- 
-

## **Q5**

Why use string.c\_str()?

Answer:

POSIX APIs use C style char\*

---

---

## **Q6**

Difference between mkdir and mkdir -p?

Answer:

mkdir → single level  
mkdir -p → recursive parent creation

---

---

## **Q7**

Is directory special?

Answer:

No. It's a file storing name → inode mappings.

---

---

## **Q8**

What happens if directory exists?

Answer:

mkdir returns -1  
errno = EEXIST

---

---

## **Q9**

Thread safety?

Answer:

mkdir is atomic at kernel level

---

---

## **Q10 (advanced)**

Which syscall used internally?

Answer:

Linux uses:

`sys_mkdir`  
`sys_mkdirat`

# rmdir command notes



```
#include "shell.hpp"
#include <iostream>
#include <sys/stat.h>
#include <unistd.h>
using namespace std;
void Shell::removeFile(const string& path){

    struct stat st;
    if (lstat(path.c_str(), &st) == -1) {
        perror("rm");
        return;
    }

    if (S_ISDIR(st.st_mode))
        rmdir(path.c_str());
    else
        unlink(path.c_str());

}
```

```
void Shell::removeFile(const string& path){
```

- 👉 Shell class ka function
  - 👉 input: file/directory ka path
  - 👉 kaam: delete karna
- 
- 

♦ **1 struct stat info;**

```
struct stat info;
```

### Kya hai?

Kernel structure jisme **file metadata store hota hai**

Contains:

```
st_mode    → file type + permissions  
st_size    → size  
st_uid     → owner  
st_gid     → group  
st_mtime   → modified time  
st_ino     → inode number
```

👉 Basically **inode ka data**

---

---

♦ **2 stat()**

```
if(stat(path.c_str(),&info)!=0){
```

**Syscall:**

```
stat()
```

**Kya karta hai?**

👉 "File ka metadata lao"

Internally:

```
path → directory lookup → inode → fill struct stat
```

---

**Return value**

```
0 → success  
-1 → fail
```

So:

```
!= 0 → error
```

---

## Fail kab hoga?

Case	errno
file not exist	ENOENT
permission denied	EACCES

---

---

- ◆ **3 perror()**

```
perror("rm");
```

## kya karta hai?

Automatically prints:

```
rm: <error message>
```

Example:

```
rm: No such file or directory
```

👉 uses errno internally

Better than manually:

```
strerror(errno)
```

---

---

- ◆ **4 return**

```
return;
```

Agar file exist hi nahi karti → delete ka sense nahi  
👉 exit function early

---

---

## 🔥 Next: file type check

---

### ◆ 5 S\_ISREG

```
if(S_ISREG(info.st_mode)){
```

**Kya hai?**

Macro

Checks:

```
Is this a regular file?
```

Internally:

```
(info.st_mode & S_IFMT) == S_IFREG
```

### Regular file examples

```
.txt  
.cpp  
.log
```

---

---

### ◆ 6 unlink()

```
if(unlink(path.c_str())!=0){
```

## Syscall

`unlink()`

### Important concept

👉 File delete **actually means removing directory entry**

Linux me:

`filename → inode`

`unlink:`

```
remove filename link  
decrement link count  
if count=0 → free inode + blocks
```

---

### Interview gold line:

👉 **unlink removes name, not immediately the data**

If file open hai:

`data stays until last fd closed`

---

---

---



## Directory case

---

- ♦  **S\_ISDIR**

```
else if(S_ISDIR(info.st_mode)){
```

Check:

👉 directory hai ya nahi

---

---

## ◆ 8 rmdir()

```
if(rmdir(path.c_str())!=0){
```

### Syscall

```
rmdir()
```

### Kya karta hai?

👉 empty directory delete

---

### Important:

```
rmdir only works if directory empty
```

Otherwise:

```
ENOTEMPTY
```

---

### Why separate from unlink?

Because:

- directories special structure hoti hain
- contains . and ..
- safety reasons

So:

```
files → unlink  
dirs → rmdir
```

---

---

## ◆ else

```
cout<<"Unsupport file type\n";
```

Other types:

- symlink
- socket
- FIFO
- device file

Abhi handle nahi kiya

---

---

## 🔥 Full flow (simple)

Program logic:

```
stat(path)  
  
if not exist → error  
  
if regular → unlink  
if directory → rmdir  
else → unsupported
```

👉 basically mini rm



## Important OS Concepts hidden here

This small code tests:

inode

stat

file types

unlink vs rmdir

errno

system calls

Interviewers love this.

---

---

## 🔥 Interview Questions (Intermediate → Advanced)

---



### Intermediate

#### Q1

Why use stat before unlink?

👉 To check file type and existence.

---

## Q2

Difference between unlink and rmdir?

- 👉 unlink → files
  - 👉 rmdir → empty directories
- 

## Q3

What happens if you unlink an open file?

- 👉 file removed from directory
  - 👉 data stays until fd closed
- 

## Q4

Why stat returns metadata not filename?

- 👉 filename stored in directory entry, not inode
- 
- 

## 🟡 Advanced

## Q5

Why stat + unlink causes race condition?

Between:

```
stat()  
unlink()
```

file might change (TOCTOU bug)

Safer:

- 👉 try unlink directly and handle error
-

## Q6

Difference between stat and lstat?

- 👉 stat follows symlink
- 👉 lstat gives symlink info

For rm:

`should use lstat`

---

otherwise symlink target delete ho sakte hai

---

## Q7

How does rm -r work internally?

- 👉 recursion:

`opendir`  
`readdir`  
`stat`  
`unlink/rmdir`

---

## Q8

What is link count?

inode field:

`st_nlink`

file deleted only when:

`link count == 0 AND no open fds`

---

## Q9

Why directories need execute permission?

Without x:

cannot cd or traverse

---

## **Q10 (very common)**

How would you improve this code?

Expected answer:

- use lstat
- handle symlinks
- recursive delete
- better error handling
- avoid TOCTOU
- use unlinkat()

# catFile command

CODE: catFile.cpp

```
#include "shell.hpp"
#include <fcntl.h>
#include <unistd.h> //for read(),write(),close()
#include <iostream>

void Shell::catFile(const string& path) {
    //open file read only
    int fd=open(path.c_str(),O_RDONLY);

    if(fd== -1){
        perror("open");
        return;
    }

    const size_t BUFFER_SIZE=4096; //linux page size=2kb
    char buffer[BUFFER_SIZE];

    ssize_t bytes;

    cout<<"=====FILE
CONTAINS=====";
    cout<<"=====END - OF -
FILE====="<<endl;
}
```

# Example

```
int fd = open("a.txt", O_RDONLY);
```

Suppose:

$$fd = 3$$

## Means:

👉 "Kernel ne file a.txt ko index 3 pe map kar diya"

**Process ke paas FD table hota hai**

## FD table (per process)

```
0 → stdin  
1 → stdout  
2 → stderr  
3 → a.txt
```

```
4 → pipe read  
5 → socket
```

👉 FD is just index into this table

### ◆ Line 1

```
int fd = open(path.c_str(), O_RDONLY);
```

### Syscall:

```
open()
```

### What happens internally?

Kernel:

1. find file inode
2. permission check
3. create open file object
4. add entry in FD table
5. return FD number

---

### Return:

```
>=0 → fd  
-1 → error
```

### ◆ Line

```
const size_t BUFFER_SIZE = 4096;
```

### Why 4096?

👉 4KB = Linux page size

Benefits:

- fewer syscalls
  - faster
  - page aligned
- 

## Important:

More read calls = slower

Bad:

read 1 byte each time ✗

Good:

read 4KB chunks ✓

---

## Interview question

Why bigger buffer improves performance?

👉 fewer syscalls → fewer kernel switches

---

---

### ♦ Line

char buffer[BUFFER\_SIZE];

Memory to store read data.

### ♦ Line

```
while((bytes = read(fd, buffer, BUFFER_SIZE)) > 0)
```

### Syscall:

```
read()
```

### What it does:

👉 copies data:

kernel → user buffer

---

### Returns:

value	meaning
e	
>0	bytes read
0	EOF
-1	error

### Example:

File size = 10KB

Flow:

```
4096  
4096  
1808  
0 (EOF)
```

### ◆ Line

```
write(1, buffer, bytes);
```

### Syscall:

```
write()
```

## What is 1??

👉 FD 1 = stdout

Default mapping:

```
0 → stdin  
1 → stdout  
2 → stderr
```

So:

```
write(1, ...) → print on screen
```

---

## Example:

```
write(1, "hello", 5)
```

prints:

```
hello
```



## Visual diagram (very important)

---

## During execution

Process FD table

```
0 → stdin  
1 → stdout  
2 → stderr  
3 → file.txt ← fd returned by open
```

Then:

```
read(3, buffer)
write(1, buffer)
```

---

---

## 🔥 Entire flow summary

```
open → get fd
loop:
    read(fd)
    write(stdout)
close(fd)
```

👉 Exactly how cat works internally

## 🔥 One-line memory tricks

Remember:

### FD basics

```
0 stdin
1 stdout
2 stderr
```

### cat logic

```
open → read → write → close
```

### performance

```
big buffer = fewer syscalls = faster
```

## Q1. What is a File Descriptor (FD)?

### Answer

File descriptor is an **integer handle** that refers to an **open file object in the kernel**.

Each process has an **FD table**:

```
0 → stdin  
1 → stdout  
2 → stderr  
3+ → files/sockets/pipes
```

`open()` returns this index.

### One-liner (interview)

👉 FD is an index into the per-process file descriptor table.

---

---

## Q2. What happens internally when `open()` is called?

### Answer (steps)

Kernel:

1. path lookup → inode
2. permission check
3. create open file description
4. add entry to FD table
5. return fd

### Diagram

```
fd → file table → inode → disk blocks
```

## One-liner

👉 open maps file to kernel object and returns fd.

---

---

## ✓ Q3. Why read()/write() instead of printf()/cout?

### Answer

read/write	printf/cout
syscall	library
unbuffered	buffered
faster, low-level	convenient

System tools use read/write for:

- performance
- control
- binary safety

## One-liner

👉 read/write are raw syscalls; printf is buffered user-space.

---

---

## ✓ Q4. Why use large buffer (4096)?

### Answer

Because:

- 4096 = page size

- fewer syscalls
- fewer kernel switches
- faster

Bad:

`read 1 byte → slow`

Good:

`read 4KB → fast`

### One-liner

👉 Larger buffer reduces syscall overhead.

---

---

## ✓ Q5. What does `read()` return and why `ssize_t`?

### Answer

`>0 → bytes read`  
`0 → EOF`  
`-1 → error`

`ssize_t` is signed so it can return -1.

### One-liner

👉 `ssize_t` allows negative error value.

---

---

## ✓ Q6. How does `cat file > out.txt` work internally?

## Answer

Shell does:

```
open(out.txt)
dup2(fd, 1)
exec(cat)
```

Now:

```
write(1,...) → goes to file
```

## Key concept

👉 stdout is just FD 1

## One-liner

👉 redirection uses dup2 to replace stdout.

---

---



## Q7. What happens if you don't close(fd)?

## Answer

FD leak:

- kernel resources not freed
- eventually:

Too many open files

Limit:

```
ulimit -n
```

## One-liner

👉 not closing causes file descriptor leak.

---

---

## ✓ Q8. Difference between text file and binary file in Linux?

### Answer

Linux:

No difference

Everything is bytes.

read/write treat both same.

Only terminal interprets ASCII.

## One-liner

👉 Linux has no text/binary distinction.

---

---

## ✓ Q9. Why is cat slow for huge files? How to optimize?

### Problem

Too many read/write syscalls.

### Solutions

- bigger buffer
- use mmap

- use `sendfile`
- async I/O

## mmap approach

file mapped directly to memory  
no copy  
faster

## One-liner

👉 use mmap or sendfile to reduce copies.

---

---

## ✓ Q10. Explain full data path of `read()`.

### Deep (very common advanced question)

When you do:

```
read(fd, buf, 4096);
```

Flow:

disk → kernel page cache → copy to user buffer

So 2 copies:

1. disk → kernel
2. kernel → user

mmap removes second copy.

## One-liner

👉 read copies kernel buffer to user space.

## 🔥 Rapid 30-sec Revision (before interview)

Remember only this:

```
FD = file handle
open → get fd
read → bytes from kernel
write → bytes to stdout(1)
4096 buffer → fewer syscalls
close → free resources
Linux files = bytes only
dup2 → redirection
mmap/sendfile → faster
```

# copyFile command

```
#include "shell.hpp"
#include <fcntl.h>
#include <unistd.h>
#include <iostream>
#include <cstring>

void Shell::copyFile(const string& src, const string& dest){
    //open file
    int src_fd=open(src.c_str(),O_RDONLY);
    if(src_fd<0){
        perror("open source");
        return;
    }

    //open/create destination file

    int dest_fd=open(dest.c_str(),
                     O_WRONLY|O_CREAT|O_TRUNC,
                     0644);

    if(dest_fd <0){
        perror("open destination");
        close(src_fd);
        return;
    }

    const int BUFFER_SIZE=4096;
    char buffer[BUFFER_SIZE];
    ssize_t bytes;

    //read -> write loop
    while((bytes=read(src_fd,buffer,BUFFER_SIZE))>0){
        if(write(dest_fd,buffer,bytes)!=bytes){

    
```

```
    perror("write");
    break;
}
}
if(bytes<0)
    perror("read");
close(src_fd);
close(dest_fd);
}
```

## ◆ STEP 1 — Open source file

```
int src_fd = open(src.c_str(), O_RDONLY);
```

---

## What happens internally (kernel steps)

When you call `open()`:

Kernel:

### 1. Path lookup

`/home/a.txt` → `inode`

### 2. Permission check

Read allowed?

### 3. Create open file object

Stores:

- offset

- flags
- inode pointer

#### 4. Add to FD table

fd → open file object

#### 5. return fd

Example:

src\_fd = 3

---

### Why int?

FD = integer index.

---

### Interview Q

#### Q: What does open return internally?

👉 index into per-process file descriptor table.

---

```
if(src_fd < 0){  
    perror("open source");  
    return;  
}
```

---

### Why check?

Because:

-1 → failure

Possible reasons:

- file not exist
  - permission denied
  - too many open files
- 

## Why perror?

Prints:

open source: No such file or directory

Uses errno.

---

## Interview Q

**Q: Why must every syscall be checked?**

👉 kernel can fail anytime → never assume success.

---

---

## ◆ STEP 2 — Open destination file

```
int dest_fd = open(dest.c_str(),  
                   O_WRONLY|O_CREAT|O_TRUNC,
```

`0644);`

---

## Flags breakdown (very important)

### `O_WRONLY`

write only

---

### `O_CREAT`

Create file if not exists

---

### `O_TRUNC`

If file exists:

`size = 0`

overwrite

---

### Combined meaning

`open for writing,`  
`create if needed,`  
`clear old content`

Exactly cp behavior.

---

---

## `0644 (permissions)`

Octal:

```
owner 6 → rw-
group 4 → r--
other 4 → r--
```

Meaning:

**rw-r--r--**

Owner can write, others read only.

Default safe permission for files.

---

## Interview Q

**Q: Why cp uses 0644 not 0755?**

👉 normal files shouldn't be executable.

---

---

```
if(dest_fd <0){
    perror("open destination");
    close(src_fd);
    return;
}
```

---

---

**Why close(src\_fd)?**

If dest open fails:

👉 avoid FD leak

VERY important.

---

## Interview Q

**Q: What happens if we don't close?**

👉 FD leak → eventually “Too many open files”.

---

---

## ◆ STEP 3 — Buffer

```
const int BUFFER_SIZE=4096;  
char buffer[BUFFER_SIZE];
```

---

## Why 4096?

Because:

4KB = Linux page size

Benefits:

- page aligned
  - optimal disk I/O
  - fewer syscalls
  - faster
- 

## Performance logic

Bad:

```
read 1 byte → 1M syscalls
```

Good:

```
read 4KB → 250 syscalls
```

---

## Interview Q

**Q: Why large buffers improve performance?**

👉 reduces syscall overhead.

---

---

## ◆ STEP 4 — read → write loop

```
while((bytes=read(src_fd,buffer,BUFFER_SIZE))>0){
```

---

## read() internals

Flow:

```
disk → kernel page cache → copy to buffer
```

Return:

```
>0 bytes  
0 EOF  
-1 error
```

---

## Why ssize\_t?

Signed because:

-1 needed for error

---

---

```
if(write(dest_fd, buffer, bytes)!=bytes){
```

---

---

## write() internals

Flow:

user buffer → kernel → disk

---

## Why check != bytes ?

Because:

write may:

- partial write
- interrupted

So:

write < bytes → incomplete

Important correctness check.

---

## Interview Q (advanced)

**Q: Can write write fewer bytes than requested?**

👉 Yes (partial write). Must loop until complete.

Most candidates miss this.

---

---

```
perror("write");  
break;
```

Stop copying on error.

---

---

- ◆ **STEP 5 — read error check**

```
if(bytes<0)  
    perror("read");
```

If read returned -1 → error.

---

---

- ◆ **STEP 6 — close**

```
close(src_fd);  
close(dest_fd);
```

---

## **close() internals**

Kernel:

- remove FD entry
  - decrement ref count
  - flush buffers
  - free resources
- 

## Why mandatory?

Otherwise:

- memory leak
  - FD leak
  - data may not flush
- 

## Interview Q

**Q: When is data actually written to disk?**

👉 usually on close or flush.

---

---

---



## FULL FLOW DIAGRAM

open(src) → fd 3  
open(dest) → fd 4

```
loop:  
    read(3) → buffer  
    write(4) → disk
```

close both

---

---



## Deep Kernel Picture (advanced)

```
disk  
↓  
page cache  
↓ copy  
user buffer  
↓ copy  
page cache  
↓  
disk
```

2 copies.

Better:

sendfile/mmap → zero copy

---

---



## High ROI Interview Questions

---



Q1

Difference between file descriptor and FILE\* ?

- 👉 fd = kernel
  - 👉 FILE\* = user buffered
- 

## 🔥 Q2

How to make copy faster?

- 👉 mmap
  - 👉 sendfile
  - 👉 splice
  - 👉 bigger buffer
- 

## 🔥 Q3

What happens if process crashes before close?

- 👉 kernel auto closes all FDs
- 

## 🔥 Q4

What if src and dest are same file?

- 👉 data corruption
  - (real cp checks inode equality)
- 

## 🔥 Q5

Why open dest before reading?

- 👉 fail fast, avoid partial copy
-

## Q6

What happens after fork?

- 👉 child inherits same FDs
- 

## Q7

How many copies happen in read/write?

- 👉 two copies
- 

## Q8

Why O\_TRUNC needed?

- 👉 clear old content
- 

## Q9

What is partial write?

- 👉 write returns less than requested
- 

## Q10

How would you implement cp -r ?

- 👉 recursion + opendir + readdir
- 
-

# One-line memory summary

```
open → get fd  
read → bytes in  
write → bytes out  
buffer 4KB  
check partial write  
close always  
0644 safe permission
```

👉 FD = integer number that represents an open file

Linux rule:

Everything is a file

So FD can represent:

- regular file
  - directory
  - pipe
  - socket
  - terminal
- 

## Example

```
int fd = open("a.txt", O_RDONLY);
```

Suppose return:

```
fd = 3
```

Means:

👉 "Kernel ne a.txt ko FD table ke index 3 pe map kar diya"

## 🔥 PART 2 — FD Table kya hota hai?

Har process ke paas hota hai:

👉 File Descriptor Table (per process)

---

### Example table

Process FD Table

FD	Points to
-----	
0	stdin (keyboard)
1	stdout (screen)
2	stderr (screen)
3	a.txt (source)
4	b.txt (destination)

👉 FD sirf index hai  
👉 actual file kernel object me hoti hai

## 🔥 PART 3 — src\_fd and dest\_fd kya represent karte hain?

Tumhara code:

```
int src_fd = open(src.c_str(), O_RDONLY);
```

```
int dest_fd = open(dest.c_str(), O_WRONLY|O_CREAT|O_TRUNC,  
0644);
```

---

## Meaning

### src\_fd

- 👉 source file handle
- 👉 read karne ke liye

Example:

```
src_fd = 3
```

---

### dest\_fd

- 👉 destination file handle
- 👉 write karne ke liye

Example:

```
dest_fd = 4
```

---

So:

```
read(3) → from a.txt  
write(4) → into b.txt
```



## PART 5 — DRY RUN (Full flow)

---

### Initial state

### a.txt

```
hello world  
(11 bytes)
```

---

---

### b.txt

```
(empty)
```

## STEP 1 — open source

```
src_fd = open("a.txt", O_RDONLY);
```

Kernel:

```
find inode  
check permission  
create open file object  
add FD entry
```

FD table:

```
0 stdin  
1 stdout  
2 stderr  
3 a.txt
```

---

---

## STEP 2 — open destination

```
dest_fd = open("b.txt", O_WRONLY|O_CREAT|O_TRUNC, 0644);
```

Kernel:

- create file if missing
- truncate to 0
- open for writing

FD table:

4 b.txt

---

---

## 🔥 Now copy loop starts

---

### STEP 3 — read()

`bytes = read(3, buffer, 4096);`

Kernel:

`disk → page cache → copy → buffer`

buffer contains:

`"hello world"`

`bytes = 11`

---

---

### STEP 4 — write()

`write(4, buffer, 11);`

Kernel:

`buffer → page cache → disk`

b.txt becomes:

`hello world`

---

---

## **STEP 5 — next read**

`read(3)`

returns:

`0 (EOF)`

Loop ends.

---

---

## **STEP 6 — close**

`close(3);  
close(4);`

Kernel:

- flush
- free FD entries



## Full Visual Flow

---

### Kernel + User space diagram

a.txt (disk)

↓

kernel page cache

↓ copy

user buffer

↓ copy

kernel page cache

↓

b.txt (disk)

---

---



## Super simple mental model

fd 3 → read

fd 4 → write

Loop:

copy bytes from 3 → 4

---

---



## Step-by-step animation

```
[read]
buffer = "hello world"
```

```
[write]
b.txt = "hello world"
```

```
[next read]
EOF
```

Done.

---

---

## 🔥 Common Interview Questions (High ROI)

---

### Q1

What is file descriptor?

👉 index into per-process FD table

---

### Q2

Why two FDs needed?

👉 one for reading, one for writing

---

### Q3

Why O\_TRUNC needed?

- 👉 clear old content before writing
- 

## Q4

What happens if O\_CREAT not used?

- 👉 open fails if file missing
- 

## Q5

Why close important?

- 👉 release resources + flush data
- 

## Q6

How many copies happen here?

- 👉 two copies (kernel↔user)
- 

## Q7

How to optimize?

- 👉 mmap/sendfile (zero copy)
- 

## Q8

Can write return partial?

👉 yes, must handle

---

## Q9

Where is file offset stored?

👉 kernel open file description

---

## Q10 (advanced)

After fork, what happens to FDs?

👉 child inherits same FD table

---

---

# 🔥 Final Memory Cheat

Remember:

```
FD = handle  
open → get fd  
read(fd1)  
write(fd2)  
close  
flags control behavior
```

# FD table



# PART 1 — Sabse pehle reality check

## Linux rule

👉 Everything is a file

Linux ke liye:

```
file  
socket  
pipe  
terminal  
keyboard  
stdout
```

sab = FILE

So OS needs **one uniform handle**

That handle =



# File Descriptor (FD)

---

---



# PART 2 — FD kya hota hai?

## Simple definition

👉 FD = small integer number

Example:

3

4

## Why integer?

Because:

`index into table`

Fast lookup.

---

---

# 🔥 PART 3 — FD table kya hota hai?

## Definition

👉 Per-process array of open files

Har process ke paas:

`FD table`

---

## Imagine like this

Process

|

+--- FD Table (array)

---

---

# 🔥 REAL MEMORY DIAGRAM

---

## Example program

```
open("a.txt");
open("b.txt");
```

---

## FD table becomes

FD	Points to
0	stdin
1	stdout
2	stderr
3	a.txt
4	b.txt

---

- 👉 FD number = index
  - 👉 inside = pointer to kernel file object
- 
- 



## Golden Rule (remember forever)

FD is NOT file  
FD is pointer to file

---

---



## PART 4 — Default FD

Every process automatically has:

0 → stdin (keyboard)  
1 → stdout (screen)  
2 → stderr (screen)

Always.

---

So:

```
printf() → write(1, ...)
```

---

---



## PART 5 — open() kya karta hai?

---

```
int fd = open("a.txt", O_RDONLY);
```

---

Kernel:

1. find inode
  2. create open file object
  3. add entry to FD table
  4. return index
- 

Example:

```
fd = 3
```

---

So:

```
read(3, ...)
```

means:

👉 read from a.txt

---

---



## PART 6 — read() ka magic

```
read(fd, buffer, 100);
```

Kernel:

```
FD table[fd]
    ↓
file object
    ↓
disk
```

Copy data to buffer.

---

---



## PART 7 — dup2 ka asli game

```
dup2(3,1);
```

Means:

```
FD 1 → same file as FD 3
```

---

### Before

```
1 → terminal
3 → file.txt
```

---

## After

```
1 → file.txt  
3 → file.txt
```

---

So:

```
printf() → goes to file
```

---

👉 This is HOW redirection works.

---

---



## PART 8 — fork() ka effect on FD table

---

```
fork();
```

Child gets:

👉 copy of FD table

---

Example:

Before:

```
Parent  
3 → a.txt
```

After:

```
Parent 3 → a.txt  
Child   3 → a.txt
```

Both share same file object.

---

So:

👉 offset shared

---

Interview favorite question 🔥

---

---



## PART 9 — exec() ka effect

---

exec:

memory replaced  
BUT FD table remains

So:

👉 redirection still works after exec

---

This is WHY:

dup2 before exec

works.

---

---



## PART 10 — FULL FLOW example

---

Command:

```
ls > out.txt
```

---

## Step 1

open:

```
fd=3
```

---

## Step 2

dup2:

```
1 → out.txt
```

---

## Step 3

exec(ls)

---

## Step 4

ls writes:

```
write(1, ...)
```

→ file

---

---

# 🔥 Final mental model

Remember:

```
FD = index  
Table = mapping  
dup2 = change mapping
```

---

---

## 🔥 NOW your 2nd question

---

### 🔥 What is `size_t` ?

---

#### Definition

`unsigned integer for sizes`

Used for:

`memory size`  
`array length`  
`buffer length`

---

#### Why `unsigned`?

Because:

`size can't be negative`

---

#### Example

```
size_t n = read(fd, buf, 4096);
```

---

## Actual type?

Depends system:

system	size
32 bit	unsigned int
64 bit	unsigned long

---

---

## Interview answer

👉 size\_t = type for sizes returned by sizeof, malloc, read lengths

---

---

## 🔥 What is pid\_t ?

---

## Definition

type for process ID

Used for:

fork()  
waitpid()  
getpid()

---

## Example

```
pid_t pid = fork();
```

---

## Why not int?

Because:

OS may change size.

So portable typedef.

---

---

## Usually:

```
typedef int pid_t;
```

But abstracted.

---

---



## Quick cheat table

---

Type	Meaning	Used for
int	general	numbers
size_t	size	memory length
ssize_t	signed size	read/write return
pid_t	process id	fork/wait
off_t	file offset	lseek

---

---



# Top Interview Questions (FD Table)

---

## Q1

What is file descriptor?

👉 index into per-process FD table

---

## Q2

Why 0,1,2 reserved?

👉 stdin, stdout, stderr

---

## Q3

What does dup2 do?

👉 duplicate FD mapping

---

## Q4

Does exec clear FD?

👉 No

---

## Q5

Does fork copy FD table?

👉 Yes (shared objects)

---

## **Q6**

How redirect stdout?

👉 `dup2(fd,1)`

---

## **Q7**

Difference between `size_t` and `int`?

👉 unsigned, platform dependent

---

## **Q8**

Why read returns `ssize_t`?

👉 can return -1 (error)

---

## **Q9**

Where is FD table stored?

👉 kernel space, per process

---

## **Q10**

What happens if not close FD?

👉 leak resources

---

---



## Final ultra simple memory rule

---

### Remember 4 lines only

```
FD = number  
table = mapping  
dup2 = change mapping  
fork copies mapping
```

**EXEC MOST IMP ENTIRE LINUX**

```
#include "shell.hpp"
#include <unistd.h>
#include <sys/wait.h>
#include <iostream>

void Shell::execCommand(const vector<string>& tokens)
{
    pid_t pid = fork();

    if(pid == 0) // child
    {
        vector<char*> argv;

        for(const auto& t : tokens)
            argv.push_back(const_cast<char*>(t.c_str()));

        argv.push_back(nullptr); // ★ VERY IMPORTANT

        execvp(argv[0], argv.data());

        perror("exec failed");
        exit(1);
    }
    else if(pid > 0) // parent
    {
        waitpid(pid, nullptr, 0);
    }
    else
    {
        perror("fork failed");
    }
}
```

## A.cpp contains

```
#include<iostream>

using namespace std;

int main() {
    cout<<"hello world";
    return 0;
}
```

### Making .exe file

```
LINUX_CML> pwd
/home/kshivam/projects/Linux_CML/src/commands
LINUX_CML> g++ a.cpp -o a
LINUX_CML>
```

### Print output

```
LINUX_CML> ./a
hello worldLINUX_CML> ./a
hello worldLINUX_CML>
```

The screenshot shows a Windows desktop environment with a terminal window and a code editor.

**Terminal Window:**

```
LINUX_CML> pwd
/home/kshivam/projects/Linux_CML/src/commands
LINUX_CML> g++ a.cpp -o a
LINUX_CML> ./a
hello worldLINUX_CML> ./a
hello worldLINUX_CML>
```

**Code Editor:**

The code editor displays the content of `a.cpp`:

```
#include<iostream>

using namespace std;

int main() {
    cout<<"hello world";
    return 0;
}
```

The terminal window also shows the command line and the output of running the executable twice.

# Remove a.cpp & a.exe file via rm

```
LINUX_CML> ./a
hello worldLINUX_CML> ./a
hello worldLINUX_CML> rm a.cpp
LINUX_CML> rm a
LINUX_CML> ls
changeDirectory.cpp
catFile.cpp
touch.cpp
redirection.cpp
copyFile.cpp
stateFile.cpp
executeExternal.cpp
.
removeFile.cpp
history.cpp
listDirectory.cpp
echo.cpp
printHelp.cpp
..
mkdir.cpp
printWorkingDirectory.cpp

LINUX_CML>
```

Ln 8, Col 2 Spaces: 4 UTF-8 LF { } C++ ⚙️ 🔍 🔔

VS Code 12:51 PM 19-02-2026

## 🔥 FINAL PROCESS DIAGRAM

Shell

```
|  
| fork  
|  
+---- Parent ---- wait ---- continue  
|  
+---- Child ---- exec(ls) ---- run ls ---- exit
```



## Kernel Level Diagram

fork → duplicate process

exec → replace memory

wait → synchronize

---

---



## Interview Ready Notes (Memorize)

---

### Shell execution pattern

fork → exec → wait

---

### fork

creates child

---

### exec

replaces program

---

### wait

prevents zombie

---

### exec never returns on success

---

`argv` must end with `NULL`

## 🔥 STEP 1 — Compile a.cpp

You type

```
LINUX_CML> g++ a.cpp -o a
```

### Step 1

```
pid_t pid = fork();
```

Kernel:

```
create child process
```

Now:

```
Parent (shell)
Child (copy of shell)
```

---

### Step 2 (child)

Build argv:

```
{"g++", "a.cpp", "-o", "a", NULL}
```

---

### Step 3

```
execvp("g++", argv)
```

Kernel:

1. search PATH

/usr/bin/g++

2. load binary

3. replace child memory

Now child becomes:

g++ process

Shell code gone.

---

## Step 4

g++ runs:

Internally:

```
open("a.cpp")
read()
compile
write("a")
```

Creates:

a (executable)

---

## Step 5

Child exits

Parent:

```
waitpid()
```

Shell resumes.

---

---



## Now STEP 2 — Run program

You type

```
LINUX_CML> ./a
```

---



## Again → mapping to YOUR code

Shell sees:

```
cmd = "./a"
```

Not builtin → goes to:

```
execCommand()
```

---

---



## FULL INTERNAL FLOW

---

**tokens**

```
{"./a"}
```

---

## **fork()**

Creates:

```
parent shell  
child copy
```

---

## **child**

```
execvp("./a")
```

Kernel:

1. open file `./a`
  2. check executable permission
  3. load ELF binary
  4. replace memory
  5. jump to main()
- 

## **Now child is YOUR PROGRAM**

```
main() of a.cpp runs
```

Example:

If a.cpp:

```
int main(){  
    cout<<"Hello\n";  
}
```

Output:

Hello

---

## child exits

Parent:

```
waitpid()
```

Shell prompt returns.

---

---



## Visual Diagram (VERY CLEAR)

---

## Compile

```
shell
  └ fork
    └ exec(g++)
      └ creates "a"
```

---

## Run

```
shell
  └ fork
    └ exec(./a)
      └ runs main()
```



## Syscalls happening (real kernel view)

When you run:

`./a`

Kernel sequence:

```
fork  
execve  
open binary  
mmap code  
set stack  
start main  
write(stdout)  
exit  
waitpid
```

## 🔥 Tier-1 (MUST KNOW – asked everywhere)

---

### ✓ Q1 — How does a shell execute `ls -l` internally?

**Answer**

1. `fork()`
2. child → `execvp("ls", argv)`
3. parent → `waitpid()`

fork creates child, exec replaces child with ls, parent waits.

---

---

### ✓ Q2 — Difference between `fork()` and `exec()`?

**Answer**

fork	exec
creates new process	replaces current process
new PID	same PID
copies memory	loads new program
fork → duplicate	
exec → replace	



### Q3 — Does exec create a new process?

#### Answer



No  
It only **replaces memory of current process**.  
PID remains same.

---

---



### Q4 — Why fork before exec?

#### Answer

If shell directly calls exec, shell itself gets replaced.

fork allows:

- child runs program
  - parent shell survives
- 
-

## Q5 — Why argv must end with NULL in exec?

### Answer

exec doesn't know argument count.

NULL marks end of arguments.

Without it → undefined behavior.

---

---

## Q6 — What does execvp do?

### Answer

- v → vector args
- p → search PATH

Searches command in PATH and executes.

---

---

## Q7 — What happens if exec succeeds?

### Answer

It never returns.

Process memory replaced.

Next instruction is new program's `main()`.

---

---

## Q8 — Why waitpid() needed?

### Answer

Prevents zombie processes.

---

---

Collects child exit status.

## Q9 — What is zombie process?

### Answer

Child finished but parent didn't wait.

---

---

Kernel keeps entry in process table.

State: Z

---

---

## Q10 — What is returned by fork()?

### Answer

return	meaning
--------	---------

0	child
---	-------

>0	parent (child pid)
----	--------------------

-1	error
----	-------

---

---

## Tier-2 (Intermediate / System design interviews)

---

## Q11 — What is copy-on-write in fork?

## Answer

Memory is NOT copied immediately.

Parent & child share pages.

Copy happens only when one writes.

Improves performance.

---

---



## Q12 — What happens to file descriptors after fork?

## Answer

Child inherits all FDs.

They point to same open file description.

Offsets are shared.

---

---



## Q13 — Why is exec fast even for large programs?

## Answer

Because it replaces memory instead of copying.

Loads only needed pages (demand paging).

---

---



## Q14 — Difference between execv and execvp?

## Answer

`execv`

`execvp`

needs full path	searches PATH
-----------------	---------------

---

---

## Q15 — Why use vector + data() for exec?

### Answer

exec needs `char**`.

vector is C++ container.

`.data()` gives raw pointer.

---

---

## Q16 — Can parent and child run simultaneously after fork?

### Answer

Yes.

Scheduler decides order.

No guaranteed sequence.

---

---

## Q17 — What happens if parent exits before child?

### Answer

Child becomes orphan.

Adopted by init/systemd.

---

---

## Q18 — Why is fork expensive?

### Answer

- page table copy
- kernel structures
- TLB flush

But mitigated by copy-on-write.

redirection command imp

```
#include "shell.hpp"
#include <unistd.h>
#include <sys/wait.h>
#include <fcntl.h>
#include <iostream>
#include <vector>
using namespace std;
void Shell::execRedirect(const vector<string>& tokens)
{
    int pos = -1;
    bool append = false;

    // find > or >>
    for(int i = 0; i < tokens.size(); ++i)
    {
        if(tokens[i] == ">" || tokens[i] == ">>")
        {
            pos = i;
            append = (tokens[i] == ">>");
            break;
        }
    }

    if(pos == -1 || pos + 1 >= tokens.size())
    {
        cout << "syntax error\n";
        return;
    }

    string filename = tokens[pos + 1]; //ls > ab.txt

    pid_t pid = fork();

    if(pid == 0)
    {
        int fd = open(
            filename.c_str(),
            append ? (O_WRONLY | O_CREAT | O_APPEND)
                   : (O_WRONLY | O_CREAT | O_TRUNC),
            0644
        );
    }
}
```

```
) ;

if(fd < 0)
{
    perror("open");
    exit(1);
}

dup2(fd, STDOUT_FILENO);
close(fd);

vector<char*> argv;

for(int i = 0; i < pos; ++i)
    argv.push_back(const_cast<char*>(tokens[i].c_str()));

argv.push_back(nullptr);

execvp(argv[0], argv.data());

perror("exec failed");
exit(1);
}
else
{
    wait(nullptr);
}
}
```

## What does **>** actually mean?

When you type:

```
ls > out.txt
```

You are NOT telling **ls** to write to file.

👉 You are telling **shell**:

"Before running ls, redirect stdout to file"

So:

shell changes FD 1 → file  
then runs ls normally

---



## Core trick (remember forever)

```
stdout = FD 1  
dup2(file_fd, 1)
```

Now anything printed → file

---

---



## Let's walk your code line-by-line

---

### ◆ Function header

```
void Shell::execRedirect(const vector<string>& tokens)
```

Input example:

```
ls > out.txt
```

tokens:

```
{"ls", ">", "out.txt"}
```



## Step 1 — find redirection operator

---

```
int pos = -1;  
bool append = false;
```

Meaning:

pos → index of > or >>  
append → whether >>

---

---

```
for(int i = 0; i < tokens.size(); ++i)
```

Loop through tokens.

---

---

```
if(tokens[i] == ">" || tokens[i] == ">>")
```

Check:

> overwrite  
>> append

---

---

```
pos = i;  
append = (tokens[i] == ">>");  
break;
```

Store:

Example:

```
tokens = {"ls", ">", "out.txt"}  
  
pos = 1  
append = false
```

---

---



## Step 2 — syntax check

```
if(pos == -1 || pos + 1 >= tokens.size())
```

Cases:

- ✗ no >
- ✗ missing filename

Examples:

```
ls >
```

---

---

```
cout << "syntax error\n";  
return;
```

Stops execution.

---

---



## Step 3 — filename

```
string filename = tokens[pos + 1];
```

Example:

```
"out.txt"
```

---

---

## 🔥 Step 4 — fork()

---

```
pid_t pid = fork();
```

### Why fork?

Same as always:

- 👉 child executes command
  - 👉 parent shell survives
- 
- 

## 🔥 Child process (IMPORTANT PART)

---

```
if(pid == 0)
```

Everything inside:

- 👉 child only
- 
- 

## 🔥 Step 5 — open file

---

```
int fd = open(  
    filename.c_str(),  
    append ? (O_WRONLY | O_CREAT | O_APPEND)  
           : (O_WRONLY | O_CREAT | O_TRUNC),
```

```
0644  
);
```

---

## Two modes

```
>  
O_TRUNC
```

clear file

---

```
>>  
O_APPEND
```

add to end

---

## Permissions

```
0644 → rw-r--r--
```

---

Now:

```
fd = 3
```

---

---

## 🔥 Step 6 — dup2 (MOST IMPORTANT LINE)

---

```
dup2(fd, STDOUT_FILENO);
```

**STDOUT\_FILENO = 1**

So:

`dup2(3, 1)`

---

## What dup2 does?

FD 1 now points to same file as FD 3

---

---



## FD table BEFORE dup2

```
0 → stdin  
1 → terminal  
2 → terminal  
3 → out.txt
```

---



## AFTER dup2

```
0 → stdin  
1 → out.txt  ★★★★  
2 → terminal  
3 → out.txt
```

---

Now:

👉 printf / cout / write(1) → file

---

---

`close(fd);`

Why?

Because:

- 1 already points to file
  - 3 not needed
- 
- 

## Step 7 — build argv

---

```
vector<char*> argv;
```

---

```
for(int i = 0; i < pos; ++i)
    argv.push_back(...)
```

Only before >.

Example:

```
{"ls"}
```

Not including > or filename.

---

---

```
argv.push_back(nullptr);
```

Mandatory for exec.

---

---

## Step 8 — exec

---

```
execvp(argv[0], argv.data());
```

Child becomes:

```
ls
```

But:

👉 stdout already redirected

So:

```
ls → prints → FD 1 → file
```

---

---

## 🔥 Step 9 — parent waits

---

```
wait(nullptr);
```

Shell waits.

---

---

## 🔥 FULL DRY RUN

---

Command:

```
ls > out.txt
```

---

## Step-by-step

---

## 1. parse

```
{"ls", ">", "out.txt"}
```

---

## 2. fork

```
parent + child
```

---

## 3. child open file

```
fd=3
```

---

## 4. dup2

```
1 → out.txt
```

---

## 5. exec ls

```
ls prints → goes to file
```

---

## 6. file content:

```
file1  
file2  
file3
```

---

---



## Visual Diagram

---

## Normal

`ls → stdout → terminal`

---

## With redirection

`ls → stdout → file`

---

---



## Kernel level flow

`fork`  
`open`  
`dup2`  
`execve`  
`write(1, ...)`

---

---



## SUPER IMPORTANT Concepts

---

**stdout = 1**

**stderr = 2**

**stdin = 0**

---

**dup2 rule**

`dup2(old, new)`

means:

`new → same file as old`

---

---



## Interview Questions (High ROI)

---

### Q1

How does `>` work internally?

👉 open file + `dup2(fd,1)` + `exec`

---

---

### Q2

Why `dup2` used?

👉 replace `stdout`

---

---

### Q3

Difference between `>` and `>>`?

👉 truncate vs append

---

---

### Q4

Why fork before redirect?

👉 so shell's stdout not changed

---

---

## Q5

What happens if you dup2 before fork?

👉 shell output redirected too (wrong)

---

---

## Q6

What is STDOUT\_FILENO?

👉 1

---

---

## Q7

What if close(fd) removed?

👉 still works but FD leak

---

---

## Q8

How implement `2>` ?

👉 `dup2(fd, STDERR_FILENO)`

---

---

## Q9

How implement < input.txt ?

👉 dup2(fd, STDIN\_FILENO)

---

---

## Q10

How implement ls > out 2>&1 ?

👉 dup2(1,2)